FITTYBIRD x KIABI

Projet General Report

See the clothes on you before you buy.

EPITECH 360° Framework | Project initiation | Brief v1.4.1

A photo-based avatar try-on for Kiabi. Web first, then one in-store pilot.

PROJECT TEAM

Name	Role	Owned area
Elie Zaylaa	Product Manager	Vision, manifesto, priorities, decision log, Gandalf hand-in
Boyu Fu	UX and Front-end Lead	Personas and interviews, flows, screens, app and mirror interface
Muhammad Rao	AI and Back-end Lead	Avatar pipeline, fit engine, architecture, techniques and risks
Hadeed Rauf	User Research and Interview Methodology Specialist	Interview method and guides, recruitment, Kiabi data requests, event tracking

Project window 28 Sep 2026 - 4 Nov 2026 | Gandalf deadline 6 Nov 2026

Product roadmap 9 Nov 2026 - 12 Feb 2027 | 6 two-week sprints

Version 2 | 4 November 2026 | Eleven checkpoints, one PDF each

This report is a reading guide, not a Gandalf checkpoint: the course requires one separate PDF per checkpoint and explicitly forbids uploading one combined file.

Table of contents

1. Executive Summary 3
2. The problem: what Kiabi's data says (Bootstrap) 4
3. The product: Fittybird x Kiabi (1a) 5
4. Team, roles and rituals (1b) 6
5. Users: six data-traced personas (2a) 7
6. Storymap: eight activities (2b) 8
7. Target architecture (3a) 9
8. Risk management (3b) 10
9. Roadmap, method and backlog (A, B, C) 11
10. Shared glossary and research annexes 12
11. Team contributions and ways of working 13
12. Open items and handover 14
13. Conclusion 15

Page numbers refresh automatically when the report is opened in Microsoft Word.

1. Executive Summary

The 360° Framework brief asks the team to turn an idea into a viable project, ready to be developed off the shelf by a dedicated team. Work is split into four periods and eleven checkpoints, reviewed like Agile sprints. Each checkpoint is uploaded to Gandalf as one separate PDF. The Word files are the sources; the PDFs in the submission folder are their exports.

This report is written so that a new reader, a reviewer or a development team can understand the whole project without opening eleven documents first. Every section names only the PDF files it refers to, as plain file names, and every one of them is in the submission folder. No external links, tool exports or personal working notes are cited.

Period	What it settles	Checkpoints
Bootstrap	Read the Kiabi data, choose a strategy, state the trade-offs	Bootstrap annexes (with 1a)
Period 1	Align the vision and form the team	1a, 1b
Period 2	Establish who the users are and share one vocabulary	2a, 2b, shared glossary
Period 3	Prove technical feasibility and face the risks	3a, 3b
Period 4	Set the launch path, method and initial backlog	A, B, C

Reading rule: a number written as a target or hypothesis is never a Kiabi measurement. Baselines come from the Kiabi dashboard; targets are pilot hypotheses; interview-based detail is deliberately absent until real participants exist.

All PDF files in the submission folder:

00_bootstrap-annexes.pdf | 1a_product-manifesto.pdf | 1b_team-and-rituals.pdf | 2a_interview-guides.pdf | 2a_participant-recruitment-plan.pdf | 2a_personas.pdf | 2a_stakeholder-mapping.pdf | 2a_user-research-strategy.pdf | 2b_storymap.pdf

3a_target-architecture.pdf | 3b_risk-management.pdf | A_roadmap.pdf | B_methodology.pdf | C_backlog-and-decision-log.pdf | C_backlog-and-user-stories.pdf | project-management.pdf | shared_glossary.pdf

Eleven of these are checkpoint uploads. Four are 2a supporting annexes that stay in the repository and are not uploaded, and two further files are kept for reference.

2. The problem: what Kiabi's data says (Bootstrap)

Kiabi's own dashboard shows a channel split that no amount of marketing can hide. Online, the fitting step is the weakest link: 13% satisfaction, graded 3/10, and 26.5% of items bought online are returned. In store, fitting reaches 75% satisfaction (8/10) and returns are only 3.7%. The two channels are strong at opposite ends of the journey: online research is 85% satisfied and payment 88%, while store research is 20% and payment 19%. A cross-channel customer is three times more profitable, yet the NFC loyalty card only counts shop entries, so the channels share no customer identity.

Measure	Online	In store
Satisfaction with fitting	13% (3/10)	75% (8/10)
Items returned	26.5%	3.7%
Average cart	67.53 EUR	29.08 EUR
Research stage satisfied	85%	20%
Payment stage satisfied	88%	19%
Traffic	8,792 unique visitors / month	1,238 store visits / month

Eight findings F1-F8 and four growth levers L1-L4 lead to decision B3: build for the web first (sprints 1-3), then pilot in one shop (sprints 4-6), rather than rolling out to all 85 French shops. Eight open data requests to Kiabi are logged with owners and due sprints; they cover return cost, return reasons, monthly orders, size charts, garment images and the privacy contact.

Referenced PDF files 00_bootstrap-annexes.pdf

3. The product: Fittybird x Kiabi (1a)

Fittybird turns a few photos into a personal avatar that wears Kiabi clothes and shows how they fit, per body zone, as tight, right or loose. The shopper takes photos after consent, the avatar is built and the photos are deleted, a garment is tried on, a size is chosen and the Kiabi cart takes over. The product explicitly is not a new checkout or payment system, not a try-on for children's clothes, not an exact body scan and not a beauty or social app.

Pilot indicator	Baseline	Target at pilot
Online item return rate on pilot articles	26.5%	23.5% (3-point fall, control group)
Online fitting grade among try-on users	3 / 10	5 / 10 or more
Avatar completion	none	70% of started captures
Time to first avatar	none	Under 3 minutes, median
Sessions with a try-on	none	40%
Size trust after a try-on	none	4 / 5
In-store reuse of the web avatar	card grade 7.3 in store	20% of pilot-shop visits

All seven targets are labelled as pilot hypotheses, not Kiabi figures. The manifesto also carries the value proposition canvas, six data-traced user types, the company challenges and decisions D1-D3 with the options that were given up. A twelve-slide presentation deck accompanies the checkpoint but is not itself a Gandalf file.

Referenced PDF files 1a_product-manifesto.pdf

4. Team, roles and rituals (1b)

Four named roles each have an owned area and one named backup, so a single absence cannot stop an area of work. The Product Manager owns vision, backlog, stakeholders, risk, planning, rituals and reporting, and the document includes an explicit not-his-job list. Hadeed Rauf's written acceptance of the fourth role is dated 4 November 2026.

Role	Owner	Backup
Product Manager	Elie Zaylaa	Hadeed Rauf
UX and Front-end Lead	Boyu Fu	Elie Zaylaa
AI and Back-end Lead	Muhammad Rao	Hadeed Rauf
User Research and Interview Methodology Specialist	Hadeed Rauf	Boyu Fu

Six rituals have a fixed cadence: sprint planning (first Monday, 1 hour), daily stand-up (10:00, 15 minutes), backlog refinement (Wednesday of week 1, 45 minutes), sprint review (last Friday, 45 minutes), retrospective (45 minutes) and a 15-minute risk review on the same Friday. Six retrospective dates run from 20 November 2026 to 12 February 2027, and the first full risk arbitration is placed in sprint 1.

Referenced PDF files 1b_team-and-rituals.pdf

5. Users: six data-traced personas (2a)

Six personas are built from the Kiabi dashboard and the seven market segments of the Kiabi data deck, not from invented demographics. The four covered segments represent 70% of buyers and 64% of turnover; decision P1 explains why Fashion Presses and Sans Limites wait until after the pilot. Decision P2 removes age, job and city from the cards: those belong to sprint 1 interviews and are not in the data. Each card carries a list of questions to validate in interviews instead.

Persona	Who	Evidence / success signal
Lea	Online shopper, unsure of size	Online fitting 13%, returns 26.5%; orders one size and keeps it
Karim	In-store shopper who dislikes shopping	Store research 20%, payment 19%; fewer trips to the shelves
Nadia	Shopper who struggles to find a size	Leaves with an item that fits, controls her photos
Lucas	Budget-conscious trend follower	Fewer returns, more items kept
Ines	Store associate with a tablet	Fewer wrong-size trips to the stockroom
Claire	Kiabi e-commerce and marketing	A readable return-rate effect on pilot articles

Section 3 gives six rules for using personas: writing stories, refining them, setting acceptance criteria, making trade-offs, reviewing the product and updating the cards. The file ends with a handover note stating that interviews have not been run. Four research annexes support this checkpoint but are not uploaded to Gandalf.

Referenced PDF files 2a_personas.pdf | 2a_user-research-strategy.pdf | 2a_stakeholder-mapping.pdf | 2a_participant-recruitment-plan.pdf | 2a_interview-guides.pdf

6. Storymap: eight activities (2b)

The story map follows the shopper through eight activities plus a Kiabi team area. Each macro-functionality carries a Must, Should or Nice class, a primary persona, the journey stage it improves, a main skill, a proposed owner and a sprint number. Section 3 is the course's job and skill prioritisation matrix, with nineteen rows. Decisions M1 and M2 define Must as what the pilot cannot run without.

Activity	Must have
1. Start and consent	Sign-in, consent before photos, delete my data
2. Create the avatar	Guided photo capture, avatar from photos
3. Browse and choose	Kiabi catalog, filters by category and size
4. Try on	Garment on the avatar, front and back views
5. Judge the fit	Size advice, fit view per body zone
6. Buy	Add to the Kiabi cart, hand off to Kiabi checkout
7. In store	Pairing QR code on the mirror, size request to the staff tablet
8. After purchase	One-question fit feedback and return reason

Section 4 explains how the screens are derived from the activities, how the avatar, garment, fit view and size objects are shared across the web, mirror, tablet and dashboard clients, and how consent, privacy and the D17/D18 pairing rules shape the opening screens. Section 6 gives design guidance for each screen and for the states that interrupt each main path. The course helper material states that prototyping is the next step after the story map, so no prototype file is required for this checkpoint. Version 2 aligns the map with D17/D18, event tracking, hosting, the privacy checklist and the Kiabi adapter.

Referenced PDF files 2b_storymap.pdf

7. Target architecture (3a)

The architecture is described in the four categories required by the course: Functional (what it does), Application (the software parts), Techniques (how it works) and Systems (where it runs). Eight business functions link to backlog IDs and name their persona. One Kiabi adapter is the only component that talks to Kiabi systems, starting on a mock feed of 20 articles. EU hosting and four environments keep real shopper photos out of test systems until the privacy review.

Open technical decision	Owner / settlement
SP-01 Build or buy the photo-to-avatar step	Sprint 1 spike, decision date 20 Nov 2026
SP-02 2D garment overlay for the try-on	Sprint 1 spike; technical backlog items
In-store identity	Pairing QR code (D17), loyalty card linked once by number (D18)
Retention and privacy	Photos deleted once the avatar exists; retention table and privacy items

Sections 5 and 6 name the players and nine end-to-end data flows: consent, avatar storage, catalog, try-on, card link, mirror pairing, size request, deletion and analytics. Version 2, dated 4 November 2026, adds the try-on renderer and store session service, real-time push with a polling fallback, technical backlog items and risks R15 to R19.

Referenced PDF files 3a_target-architecture.pdf

8. Risk management (3b)

Nineteen risks are brainstormed and classified as Business, People, Technique or Timing, then scored by likelihood multiplied by impact on a 1-3 scale. The heat map counts twelve red and seven amber risks. Every row has a trigger, an opening score, a response, a mitigation, a class, an owner, a helper and an empty arbitrated-score column. The version-2 classification gives three Owned, one Accepted and fifteen Managed risks.

Top risk	Answer
R6 / R7 No deep 3D or vision experience; poor photos	Sprint 1 build-or-buy spike, photo quality check guiding the shopper
R9 Body photos are sensitive	Privacy checklist before any outside photo, full review before the pilot
R18 / R19 Enabler load and single-person dependency	One Should story moves to Later; a back-end backup is named
R2 / R13 / R15 Missing Kiabi data	Eight data requests tracked, gaps written into the register

A blank arbitration session record with four member vote columns and a fifteen-row mitigation action table are ready for the first sessions proposed for 12 and 18 November 2026, with the 11 November French holiday noted. The votes are intentionally not filled in: they require the four people in the room.

Referenced PDF files 3b_risk-management.pdf

9. Roadmap, method and backlog (A, B, C)

Only three months are planned: twelve working weeks from 9 November 2026 to 12 February 2027, six two-week sprints with a two-week holiday break. The riskiest and most valuable items come first. The plan uses 112 of 120 person-days in sprints 1-5, leaving eight days of slack, under a stated assumption of three available days per member per week. Four change-course thresholds say when the plan would be revised.

Sprint	Dates	Victory to celebrate
S1	9 - 20 Nov 2026	Build-or-buy decided
S2	23 Nov - 4 Dec 2026	Everyone has their own avatar
S3	7 - 18 Dec 2026	A garment on a real avatar
S4	4 - 15 Jan 2027	A size request reaches the staff tablet
S5	18 - 29 Jan 2027	MVP ready for the pilot
S6	1 - 12 Feb 2027	Real shoppers use it

Two-week Scrum sprints run with a Kanban board that limits work in progress to four stories; every story names one persona and every review replays one persona's journey. The backlog holds 33 items, including eight enabler epics, and section 2 of C gives nineteen development-ready user stories with Given / When / Then criteria. The decision log holds D1-D22 plus B3, M1, M2, P1, P2, T1 and T2; old entries are never edited, new ones are appended.

Referenced PDF files A_roadmap.pdf | B_methodology.pdf | C_backlog-and-decision-log.pdf | C_backlog-and-user-stories.pdf

10. Shared glossary and research annexes

The shared glossary gives one meaning per term across three groups: product, Kiabi and retail, and project and method. Version 2 adds pairing QR code, store session, loyalty-card link, Kiabi adapter, try-on renderer, mock feed, retention, enabler and risk class, each with a do-not-confuse-with note. It is what keeps eleven checkpoint files using the same words for avatar, try-on, fit view and pilot.

The research workstream produced four documents: the strategy, the stakeholder map, the participant recruitment plan and four field-use interview guides with probing banks, forbidden questions and bilingual scripts. All four were exported as 2a supporting annexes and are kept in the repository but are not uploaded to Gandalf. Outputs 05 to 09 do not exist, because no interviews have taken place.

Referenced PDF files shared_glossary.pdf | 2a_interview-guides.pdf | 2a_user-research-strategy.pdf | 2a_stakeholder-mapping.pdf | 2a_participant-recruitment-plan.pdf

11. Team contributions and ways of working

Four named roles share the initiation work. Each role has an owned area and one named backup, so a single absence cannot stop an area. Ownership below is by deliverable: the day-to-day work behind each checkpoint was split across the team, and no checkpoint was produced by one member alone.

Member	Role	Main contributions	Backup
Elie Zaylaa	Product Manager	Vision and boundaries, backlog and priorities, contact with Kiabi and the school, planning and the risk register; the manifesto, team and rituals, decision log and glossary	Hadeed Rauf
Boyu Fu	UX and Front-end Lead	Personas and interviews, flows, screens, the app and mirror interface; the two journey maps, six persona cards, screen and state design, seven figures and the first user stories	Elie Zaylaa
Muhammad Rao	AI and Back-end Lead	Avatar pipeline, fit engine, API, data storage and privacy; the four-layer architecture, technical risks, capacity plan and technical backlog	Hadeed Rauf
Hadeed Rauf	User Research and Interview Methodology Specialist	Interview method and guides, buying-pattern analysis, Kiabi data requests, event tracking and KPI reporting, test plan and pilot logistics; the four research annexes	Boyu Fu

The team works in two-week Scrum sprints with a Kanban board that limits work in progress to four stories. Six rituals have a fixed cadence: sprint planning on the first Monday, a daily stand-up at 10:00, backlog refinement in week one, a sprint review on the last Friday, a retrospective, and a 15-minute risk review on the same Friday. Every story names one persona and every review replays one persona's journey end to end.

Decisions are append-only. A change is recorded as a new dated entry in the decision log; old entries are never edited. Six retrospective dates run from 20 November 2026 to 12 February 2027, and the first full risk arbitration is placed in sprint 1. The initiation phase closed on 4 November 2026, two days before the Gandalf deadline.

Referenced PDF files 1b_team-and-rituals.pdf | B_methodology.pdf | C_backlog-and-decision-log.pdf | project-management.pdf

12. Open items and handover

This section covers the whole team, not one role. The repository is a completed initiation package, but several items cannot be closed by writing another file: they need real participants, a meeting with all four members, an answer from Kiabi, or a format decision from the school. Each item below names its owner and the exact reason it remains open.

Item	Owner	What remains
OPEN.6 Sprint 1 exploratory interviews	Boyu Fu	No participants, no recordings and no verbatims; research outputs 05 to 09 are intentionally blank. Real interviewees are needed, using Hadeed Rauf's guides.
QA.5 Five missing screens and states	Boyu Fu	Design guidance is written in section 6 of the 2b story map; graphics belong to the course's next step, prototyping.
QA.4 Figma rebuild	Boyu Fu	Conditional. A source file is only needed if the school or Kiabi asks for that format.
13 core screens	Boyu Fu	Sections 4 and 6 of 2b set out how the screens and states are designed. No checkpoint PDF is required for them.
OPEN.3 Eight OPS-01 data requests	Hadeed Rauf	Return cost, return reasons, size charts, garment images and the privacy contact are requested from Kiabi; answers and their gaps are not yet available.
OPEN.12 Research outputs 05 to 09	Hadeed Rauf	Interview records, analysis, empathy maps, journey maps and a findings report must be drafted from real evidence, not assumed answers.
OPEN.15 / OPEN.16 Technical handover	Muhammad Rao	Seven open technical decisions, the two sprint 1 spikes, a named backup for the avatar pipeline and the capacity figures need dated owners and short briefs.
QA.3 Capacity figures	Muhammad Rao	One load series must be quoted consistently across the decision log, the backlog and the roadmap.
OPEN.7 Proposed decisions	Elie Zaylaa	Each proposed decision needs a dated vote or a new entry that supersedes it.
OPEN.17 Gandalf hand-in	Elie Zaylaa	Eleven checkpoint PDFs must be uploaded separately; one combined file is not acceptable.
QA.7 / OPEN.18 Contribution log and rehearsal	Elie Zaylaa	The contribution record referenced by the repository guide does not exist, and the final defence rehearsal with its Q&A still needs the team.
OPEN.14 Source files	Elie Zaylaa	The Kiabi brief and data deck sit at the workspace root; the source notes in the repository guide need to point there.
OPEN.2 Risk arbitration	All four members	Nineteen arbitrated scores and four vote columns need the team in a session; the template is deliberately blank.
OPEN.5 Weekly availability	All four members	The three-days-a-week capacity assumption needs written confirmation or a revised plan.

Two items were closed during the final review without creating a new checkpoint file. The 1a presentation deck now describes the mirror pairing QR code and the one-time loyalty-card link (D17, D18) instead of a card tap, and the 2b story map now explains how the screens and their states are designed instead of pointing to an external canvas. Neither change required a new checkpoint PDF.

Referenced PDF files project-management.pdf

13. Conclusion

Fittybird x Kiabi answers one measurable problem: Kiabi's online fitting step satisfies 13% of shoppers and drives 26.5% returns, while the same step in store satisfies 75%. The proposed answer is a photo-based avatar that shows fit per body zone, delivered on the web first and piloted in one shop, with no new checkout, no children's try-on and no claim of an exact body scan.

Across four periods the team produced eleven checkpoint documents: a data diagnosis and journey maps, a product manifesto with boundaries and seven quantified hypotheses, a named team with rituals, six data-traced personas, an eight-activity story map with job and skill priorities, a four-layer target architecture with nine data flows, a nineteen-risk register, a three-month roadmap, a Scrum and Kanban method, a 33-item backlog with decision-ready user stories, and a shared glossary. Each has its own PDF in the submission folder.

The package is ready to be handed to a development team, which is exactly what the brief asks for. Two technical spikes, the first risk arbitration, Kiabi's data answers and the sprint 1 interviews remain genuinely open, and section 12 names an owner for each. The most important of these is the interview work: the persona cards keep their questions to validate in interviews precisely so that real evidence, not invented detail, will fill them.

A reader who needs only the essentials should read sections 2, 3, 9 and 12. A reviewer checking compliance should start with section 1 and the eleven named PDFs. A developer taking over should start with sections 6, 7, 9 and 11.